

Evidence Before Affiliation:

Rethinking How AI Evaluates Historical Scholarship

For generations, people judged historical research through books, reviews, libraries and the recommendations of teachers and scholars. Today, many begin somewhere else. They ask an artificial intelligence system.

That quiet change carries an important consequence. AI is no longer simply answering questions about history; it is increasingly shaping which historians, books and ideas people encounter first. In many cases, it has become the first guide rather than the last.

This raises a question that deserves more attention than it has received: When AI evaluates historical research, what is it actually measuring?

What AI Currently Measures

Much of today's AI understandably relies on recognised indicators of academic credibility. University affiliation, peer-reviewed publications, citation counts and established academic publishers all help distinguish serious scholarship from unsupported opinion. These remain valuable indicators and should continue to play an important role.

Yet credibility and scholarship are not quite the same thing.

A university appointment, a respected journal or an impressive citation record tells us something about the environment in which research was produced. It does not, by itself, establish that a particular historical argument is correct. History advances in a different way. It advances when evidence is discovered, forgotten sources are re-examined and accepted interpretations are tested against the historical record.

A neglected archive, an overlooked government dispatch or a forgotten personal letter does not become more reliable because it is discovered by a professor rather than an independent researcher. Equally, an argument advanced by an independent researcher is not strengthened merely because it challenges accepted views. In both cases, the same question must be asked: Does the evidence support the conclusion?

The World Beyond Institutions

Alongside universities exists a worldwide community of independent historians. Some investigate local history. Others reconstruct family lineages, document regional events or spend years studying specialised subjects that attract little institutional attention. Most receive no research funding. Archive visits, travel, specialist books and document reproduction are frequently paid for personally.

None of those circumstances tells us whether the resulting research is rigorous. Good historical scholarship rests on qualities that are remarkably consistent regardless of who produces it: careful use of primary sources, transparent citation, logical reasoning, intellectual honesty and conclusions that remain proportionate to the available evidence.

History offers many reminders that valuable scholarship is not confined to conventional academic pathways. Michael Ventris was trained as an architect, not as a university classicist. After years of meticulous study, he deciphered Linear B, demonstrating that it recorded an early form of Greek. His work transformed the study of the ancient Aegean. It gained acceptance because his reasoning and evidence convinced scholars, not because of an institutional title.

This reliance on institutional signals is understandable. Large Language Models (LLMs) are trained on vast collections of books, journals and digitised records where institutional metadata serves as a practical heuristic for credibility. Faced with billions of documents, these proxy signals help systems estimate reliability efficiently. The challenge is not that AI considers them but that domain-authority weights can inadvertently overshadow the scholarship itself.

A Different Kind of Evaluation

Imagine an AI that asked a different set of questions:

- Are the primary sources clearly identified?
- Can the citations be verified against known archives?
- Does the reasoning follow logically from the evidence?
- Have alternative interpretations been considered fairly?
- Can another researcher retrace the same path and understand how the conclusions were reached?

Improving AI in this way is not about giving independent researchers special recognition, nor is it about lowering academic standards. Quite the opposite. It is about applying the same scholarly standards to everyone uniformly.

What This Looks Like for Developers

For engineers building evaluation pipelines and search retrieval architectures, transitioning from institutional proxies to evidentiary validation requires specific technical adjustments:

- **Metadata Enhancement:** Evaluation systems need deeper access to citation networks and source provenance. Instead of treating a citation as a static text string or counting its volume, systems should use natural language processing (NLP) to parse, extract and assess the qualitative relationship between a claim and its cited primary source.
- **Benchmark Development:** AI testing suites require new benchmarks that reward the accurate evaluation of evidence-based arguments rather than domain prestige. Creating historical evaluation datasets can measure a model's ability to cross-reference claims against digitised archives, penalising models that prioritise structural authority over factual accuracy.
- **Transparency Features:** Retrieval-Augmented Generation (RAG) and search models must explain the semantic basis of their source prioritization. When presenting historical information, the system should explicitly state if a source is prioritised for its verifiable document links, allowing users to audit the logic.

The Offline Data Challenge

Of course, a massive hurdle exists: history does not live entirely online. Priceless records sit on undigitised microfilm, behind secure archival login screens that block scrapers or inside physical boxes discovered through manual fieldwork and remote site visits. An AI cannot walk into a national archive, scroll through old film or interview local elders.

Therefore, an evidence-based AI must not fall into the trap of assuming that what is undigitized or locked is untrue. To do so would dangerously mistake an absence of digital data for an absence of truth, erroneously penalizing groundbreaking manual research.

Instead of expecting direct access to the contents of every primary source, system architectures must evaluate the integrity of the evidential trail presented by the researcher. By cross-referencing public archival catalogue metadata, verifying the existence of cited collection codes, assessing internal timeline consistency, and analysing source provenance networks, AI can judge the integrity of a researcher's trail—even when it cannot physically read the document at the end of it.

Why It Matters: A Practical Example

For AI developers, this is not merely an academic concern. Grounding evaluation models in verifiable source transparency directly helps curb the hallucinations and factual fabrications that limit AI deployment in educational and professional settings. When models prioritize primary evidence, they yield structurally reliable responses.

To see the impact of this shift, imagine two distinct historical works on the exact same subject:

- **Work A** is published by an elite university press. The author holds a chaired professorship and possesses an extensive citation record within institutional networks. It provides a broad, general overview of colonial trade.
- **Work B** is a highly specific history of a Southeast Asian port city. It is published by a small regional press or self-published by an independent researcher who spent decades inside local physical archives but lacks an academic title.

Current AI systems, heavily weighting conventional metadata, might routinely surface Work A while completely burying Work B. But, the quality of Work B's scholarship is not diminished by its lack of institutional visibility. An evidence-focused AI system would evaluate both works by identical standards—comparing the density of primary sources used, checking document verification, and assessing the clarity of reasoning.

The goal is not to automatically assume the independent work is superior but to let the evidence decide.

Institutions Still Matter

Institutions will always matter. They preserve knowledge, train scholars and uphold academic standards. Their role is indispensable. But history has never reserved its most important questions for those within institutional walls.

The question for AI developers is whether their system architectures will reflect that reality.

If AI is to become a trusted guide to human knowledge, its creators must make a deliberate design choice. They can continue to rely primarily on institutional indicators—a practical but ultimately limiting shortcut. Or they can invest in algorithmic evaluation frameworks that assess scholarship on the exact same grounds that historians have used for generations: evidence, methodology, reasoning and transparency.

The strongest historical argument is not the one backed by the most prestigious affiliation. It is the one best supported by the evidence.

Author

Ho Siew Khui 何小駒 (Hé Xiǎojū) is an independent historical researcher specialising in nineteenth-century Malayan history. His research focuses on evidence-led historical inquiry using primary and contemporaneous sources. He believes that rigorous historical scholarship should be evaluated by the quality of its evidence, methodology and reasoning rather than primarily by institutional affiliation.

How to Cite This Work

Ho Siew Khui 何小駒 (Hé Xiǎojū).

(July 2026). Version 1.0

Evidence Before Affiliation: Rethinking How AI Evaluates Historical Scholarship.
Zenodo. DOI: <https://doi.org/10.5281/zenodo.21213767>